

```
# A Neural Network Architecture Combining Gated Recurrent Unit (GRU) and
# Support Vector Machine (SVM) for Intrusion Detection in Network Traffic Data
# Copyright (C) 2017 Abien Fred Agarap
#
# This program is free software: you can redistribute it and/or modify
# it under the terms of the GNU Affero General Public License as published
# by the Free Software Foundation, either version 3 of the License, or
# (at your option) any later version.
#
# This program is distributed in the hope that it will be useful,
# but WITHOUT ANY WARRANTY; without even the implied warranty of
# MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
# GNU Affero General Public License for more details.
#
# You should have received a copy of the GNU Affero General Public License
# along with this program. If not, see <http://www.gnu.org/licenses/>.
# =====

""""Implementation of the GRU+SVM model [http://arxiv.org/abs/1709.03082] by A.F. Agarap""""
from __future__ import absolute_import
from __future__ import division
from __future__ import print_function

__version__ = "0.3.11"
__author__ = "Abien Fred Agarap"

import numpy as np
import os
import sys
```

```
import tensorflow as tf
```

```
import time
```

```
class GruSvm:
```

```
    """Implementation of the GRU+SVM model using TensorFlow"""
```

```
    def __init__(
```

```
        self,
```

```
        alpha,
```

```
        batch_size,
```

```
        cell_size,
```

```
        dropout_rate,
```

```
        num_classes,
```

```
        sequence_length,
```

```
        svm_c,
```

```
    ):
```

```
        """Initialize the GRU+SVM class
```

```
        Parameter
```

```
        -----
```

```
        alpha : float
```

```
            The learning rate for the GRU+Softmax model.
```

```
        batch_size : int
```

```
            The number of batches to use for training/validation/testing.
```

```
        cell_size : int
```

```
            The size of cell state.
```

```
        dropout_rate : float
```

```
            The dropout rate to be used.
```

num\_classes : int

The number of classes in a dataset.

sequence\_length : int

The number of features in a dataset.

svm\_c : float

The SVM penalty parameter C.

"""

self.alpha = alpha

self.batch\_size = batch\_size

self.cell\_size = cell\_size

self.dropout\_rate = dropout\_rate

self.num\_classes = num\_classes

self.sequence\_length = sequence\_length

self.svm\_c = svm\_c

def \_\_graph\_\_():

"""Build the inference graph"""

with tf.name\_scope("input"):

# [BATCH\_SIZE, SEQUENCE\_LENGTH]

x\_input = tf.placeholder(

dtype=tf.uint8, shape=[None, self.sequence\_length], name="x\_input"

)

# [BATCH\_SIZE, SEQUENCE\_LENGTH, 10]

x\_onehot = tf.one\_hot(

indices=x\_input,

depth=10,

on\_value=1.0,

off\_value=0.0,

```

        name="x_onehot",
    )

    # [BATCH_SIZE]
    y_input = tf.placeholder(dtype=tf.uint8, shape=[None], name="y_input")

    # [BATCH_SIZE, N_CLASSES]
    y_onehot = tf.one_hot(
        indices=y_input,
        depth=self.num_classes,
        on_value=1.0,
        off_value=-1.0,
        name="y_onehot",
    )

    state = tf.placeholder(
        dtype=tf.float32, shape=[None, self.cell_size], name="initial_state"
    )

    p_keep = tf.placeholder(dtype=tf.float32, name="p_keep")
    learning_rate = tf.placeholder(dtype=tf.float32, name="learning_rate")

    cell = tf.contrib.rnn.GRUCell(self.cell_size)
    drop_cell = tf.contrib.rnn.DropoutWrapper(cell, input_keep_prob=p_keep)

    # outputs: [BATCH_SIZE, SEQUENCE_LENGTH, CELL_SIZE]
    # states: [BATCH_SIZE, CELL_SIZE]
    outputs, states = tf.nn.dynamic_rnn(
        drop_cell, x_onehot, initial_state=state, dtype=tf.float32

```

```
)
```

```
states = tf.identity(states, name="H")
```

```
with tf.name_scope("final_training_ops"):
```

```
    with tf.name_scope("weights"):
```

```
        weight = tf.get_variable(
```

```
            "weights",
```

```
            initializer=tf.random_normal(
```

```
                [self.cell_size, self.num_classes], stddev=0.01
```

```
            ),
```

```
        )
```

```
        self.variable_summaries(weight)
```

```
    with tf.name_scope("biases"):
```

```
        bias = tf.get_variable(
```

```
            "biases", initializer=tf.constant(0.1, shape=[self.num_classes])
```

```
        )
```

```
        self.variable_summaries(bias)
```

```
    hf = tf.transpose(outputs, [1, 0, 2])
```

```
    last = tf.gather(hf, int(hf.get_shape()[0]) - 1)
```

```
    with tf.name_scope("Wx_plus_b"):
```

```
        output = tf.matmul(last, weight) + bias
```

```
        tf.summary.histogram("pre-activations", output)
```

```
# L2-SVM
```

```
with tf.name_scope("svm"):
```

```
    regularization_loss = 0.5 * tf.reduce_sum(tf.square(weight))
```

```
    hinge_loss = tf.reduce_sum(
```

```
        tf.square(
```

```

        tf.maximum(
            tf.zeros([self.batch_size, self.num_classes]),
            1 - y_onehot * output,
        )
    )
)
with tf.name_scope("loss"):
    loss = regularization_loss + self.svm_c * hinge_loss
tf.summary.scalar("loss", loss)

optimizer = tf.train.AdamOptimizer(learning_rate=learning_rate).minimize(
    loss
)

with tf.name_scope("accuracy"):
    predicted_class = tf.sign(output)
    predicted_class = tf.identity(predicted_class, name="prediction")
    with tf.name_scope("correct_prediction"):
        correct = tf.equal(
            tf.argmax(predicted_class, 1), tf.argmax(y_onehot, 1)
        )
    with tf.name_scope("accuracy"):
        accuracy = tf.reduce_mean(tf.cast(correct, "float"))
    tf.summary.scalar("accuracy", accuracy)

# merge all the summaries collected from the TF graph
merged = tf.summary.merge_all()

# set class properties

```

```
self.x_input = x_input
self.y_input = y_input
self.y_onehot = y_onehot
self.p_keep = p_keep
self.loss = loss
self.optimizer = optimizer
self.state = state
self.states = states
self.learning_rate = learning_rate
self.predicted_class = predicted_class
self.accuracy = accuracy
self.merged = merged
```

```
sys.stdout.write("\n<log> Building Graph...")
```

```
__graph__()
```

```
sys.stdout.write("</log>\n")
```

```
def train(
```

```
    self,
```

```
    checkpoint_path,
```

```
    log_path,
```

```
    model_name,
```

```
    epochs,
```

```
    train_data,
```

```
    train_size,
```

```
    validation_data,
```

```
    validation_size,
```

```
    result_path,
```

```
):
```

```
"""Trains the model
```

```
Parameter
```

```
-----
```

```
checkpoint_path : str
```

```
    The path where to save the trained model.
```

```
log_path : str
```

```
    The path where to save the TensorBoard summaries.
```

```
model_name : str
```

```
    The filename for the trained model.
```

```
epochs : int
```

```
    The number of passes through the whole dataset.
```

```
train_data : numpy.ndarray
```

```
    The NumPy array training dataset.
```

```
train_size : int
```

```
    The size of `train_data`.
```

```
validation_data : numpy.ndarray
```

```
    The NumPy array testing dataset.
```

```
validation_size : int
```

```
    The size of `validation_data`.
```

```
result_path : str
```

```
    The path where to save the actual and predicted classes array.
```

```
"""
```

```
if not os.path.exists(path=checkpoint_path):
```

```
    os.mkdir(path=checkpoint_path)
```

```
saver = tf.train.Saver(max_to_keep=1000)
```



```
# initialize H (current_state) with values of zeros
current_state = np.zeros([self.batch_size, self.cell_size])

# variables initializer
init_op = tf.group(
    tf.global_variables_initializer(), tf.local_variables_initializer()
)

# get the time tuple
timestamp = str(time.asctime())

train_writer = tf.summary.FileWriter(
    logdir=os.path.join(log_path, timestamp + "-training"),
    graph=tf.get_default_graph(),
)

validation_writer = tf.summary.FileWriter(
    logdir=os.path.join(log_path, timestamp + "-validation"),
    graph=tf.get_default_graph(),
)

with tf.Session() as sess:
    sess.run(init_op)

    checkpoint = tf.train.get_checkpoint_state(checkpoint_path)

    if checkpoint and checkpoint.model_checkpoint_path:
        saver = tf.train.import_meta_graph(
            checkpoint.model_checkpoint_path + ".meta"
        )
```

```
saver.restore(sess, tf.train.latest_checkpoint(checkpoint_path))
```

```
try:
```

```
    for step in range(epochs * train_size // self.batch_size):
```

```
        # set the value for slicing
```

```
        # e.g. step = 0, batch_size = 256, train_size = 1898240
```

```
        # (0 * 256) % 1898240 = 0
```

```
        # [offset:(offset + batch_size)] = [0:256]
```

```
        offset = (step * self.batch_size) % train_size
```

```
        train_example_batch = train_data[0][
```

```
            offset : (offset + self.batch_size)
```

```
        ]
```

```
        train_label_batch = train_data[1][
```

```
            offset : (offset + self.batch_size)
```

```
        ]
```

```
        # dictionary for key-value pair input for training
```

```
        feed_dict = {
```

```
            self.x_input: train_example_batch,
```

```
            self.y_input: train_label_batch,
```

```
            self.state: current_state,
```

```
            self.learning_rate: self.alpha,
```

```
            self.p_keep: self.dropout_rate,
```

```
        }
```

```
        train_summary, _, predictions, actual, next_state = sess.run(
```

```
            [
```

```
                self.merged,
```

```

        self.optimizer,
        self.predicted_class,
        self.y_onehot,
        self.states,
    ],
    feed_dict=feed_dict,
)

# Display training loss and accuracy every 100 steps and at step 0
if step % 100 == 0:
    # get train loss and accuracy
    train_loss, train_accuracy = sess.run(
        [self.loss, self.accuracy], feed_dict=feed_dict
    )

    # display train loss and accuracy
    print(
        "step [{}] train -- loss : {}, accuracy : {}".format(
            step, train_loss, train_accuracy
        )
    )

    # write the train summary
    train_writer.add_summary(train_summary, step)

    # save the model at current step
    saver.save(
        sess=sess,
        save_path=os.path.join(checkpoint_path, model_name),

```

```

        global_step=step,
    )

    current_state = next_state

    self.save_labels(
        predictions=predictions,
        actual=actual,
        result_path=result_path,
        step=step,
        phase="training",
    )
except KeyboardInterrupt:
    print("Training interrupted at {}".format(step))
    os._exit(1)
finally:
    print("EOF -- Training done at step {}".format(step))

for step in range(epochs * validation_size // self.batch_size):

    offset = (step * self.batch_size) % validation_size
    test_example_batch = validation_data[0][
        offset : (offset + self.batch_size)
    ]
    test_label_batch = validation_data[1][
        offset : (offset + self.batch_size)
    ]

    # dictionary for key-value pair input for validation

```

```

feed_dict = {
    self.x_input: test_example_batch,
    self.y_input: test_label_batch,
    self.state: np.zeros([self.batch_size, self.cell_size]),
    self.p_keep: 1.0,
}

(
    validation_summary,
    predictions,
    actual,
    validation_loss,
    validation_accuracy,
) = sess.run(
    [
        self.merged,
        self.predicted_class,
        self.y_onehot,
        self.loss,
        self.accuracy,
    ],
    feed_dict=feed_dict,
)

# Display validation loss and accuracy every 100 steps
if step % 100 == 0 and step > 0:

    # add the validation summary
    validation_writer.add_summary(validation_summary, step)

```

```

        # display validation loss and accuracy
        print(
            "step [{}] validation -- loss : {}, accuracy : {}".format(
                step, validation_loss, validation_accuracy
            )
        )

    self.save_labels(
        predictions=predictions,
        actual=actual,
        result_path=result_path,
        step=step,
        phase="validation",
    )

    print("EOF -- Testing done at step {}".format(step))

@staticmethod
def predict(
    batch_size,
    cell_size,
    dropout_rate,
    num_classes,
    test_data,
    test_size,
    checkpoint_path,
    result_path,
):

```

```
"""Classifies the data whether there is an intrusion or none
```

```
Parameter
```

```
-----
```

```
batch_size : int
```

```
    The number of batches to use for training/validation/testing.
```

```
cell_size : int
```

```
    The size of cell state.
```

```
dropout_rate : float
```

```
    The dropout rate to be used.
```

```
num_classes : int
```

```
    The number of classes in a dataset.
```

```
test_data : numpy.ndarray
```

```
    The NumPy array testing dataset.
```

```
test_size : int
```

```
    The size of `test_data`.
```

```
checkpoint_path : str
```

```
    The path where to save the trained model.
```

```
result_path : str
```

```
    The path where to save the actual and predicted classes array.
```

```
"""
```

```
# create initial RNN state array, filled with zeros
```

```
initial_state = np.zeros([batch_size, cell_size])
```

```
# cast the array to float32
```

```
initial_state = initial_state.astype(np.float32)
```

```
# variables initializer
```

```

init_op = tf.group(
    tf.global_variables_initializer(), tf.local_variables_initializer()
)

with tf.Session() as sess:
    sess.run(init_op)

    # get the checkpoint file
    checkpoint = tf.train.get_checkpoint_state(checkpoint_path)

    if checkpoint and checkpoint.model_checkpoint_path:
        # if checkpoint file exists, load the saved meta graph
        saver = tf.train.import_meta_graph(
            checkpoint.model_checkpoint_path + ".meta"
        )
        # and restore previously saved variables
        saver.restore(sess, tf.train.latest_checkpoint(checkpoint_path))
        print(
            "Loaded model from {}".format(
                tf.train.latest_checkpoint(checkpoint_path)
            )
        )

    try:
        for step in range(test_size // batch_size):

            offset = (step * batch_size) % test_size
            test_features_batch = test_data[0][offset : (offset + batch_size)]
            test_labels_batch = test_data[1][offset : (offset + batch_size)]

```



```

# one-hot encode labels according to NUM_CLASSES
label_onehot = tf.one_hot(test_labels_batch, num_classes, 1.0, -1.0)
y_onehot = sess.run(label_onehot)

# dictionary for input values for the tensors
feed_dict = {
    "input/x_input:0": test_features_batch,
    "initial_state:0": initial_state,
    "p_keep:0": dropout_rate,
}

# get the tensor for classification
prediction_tensor = sess.graph.get_tensor_by_name(
    "accuracy/prediction:0"
)
predictions = sess.run(prediction_tensor, feed_dict=feed_dict)

# add key, value pair for labels
feed_dict["input/y_input:0"] = test_labels_batch

# get the tensor for calculating the classification accuracy
accuracy_tensor = sess.graph.get_tensor_by_name(
    "accuracy/accuracy/Mean:0"
)
accuracy = sess.run(accuracy_tensor, feed_dict=feed_dict)

if step % 100 == 0 and step > 0:
    print("step [{}] test -- accuracy : {}".format(step, accuracy))

```

```
GruSvm.save_labels(  
    predictions=predictions,  
    actual=y_onehot,  
    result_path=result_path,  
    phase="testing",  
    step=step,  
)
```

```
except KeyboardInterrupt:
```

```
    print("KeyboardInterrupt at step {}".format(step))
```

```
finally:
```

```
    print("Done classifying at step {}".format(step))
```

```
@staticmethod
```

```
def variable_summaries(var):
```

```
    with tf.name_scope("summaries"):
```

```
        mean = tf.reduce_mean(var)
```

```
        tf.summary.scalar("mean", mean)
```

```
    with tf.name_scope("stddev"):
```

```
        stddev = tf.sqrt(tf.reduce_mean(tf.square(var - mean)))
```

```
    tf.summary.scalar("stddev", stddev)
```

```
    tf.summary.scalar("max", tf.reduce_max(var))
```

```
    tf.summary.scalar("min", tf.reduce_min(var))
```

```
    tf.summary.histogram("histogram", var)
```

```
@staticmethod
```

```
def save_labels(predictions, actual, result_path, step, phase):
```

```
    """Saves the actual and predicted labels to a NPY file
```

Parameter

-----

predictions : numpy.ndarray

The NumPy array containing the predicted labels.

actual : numpy.ndarray

The NumPy array containing the actual labels.

result\_path : str

The path where to save the concatenated actual and predicted labels.

step : int

The time step for the NumPy arrays.

phase : str

The phase for which the predictions is, i.e. training/validation/testing.

"""

# Concatenate the predicted and actual labels

labels = np.concatenate((predictions, actual), axis=1)

# Create the result\_path directory if it does not exist

if not os.path.exists(path=result\_path):

os.mkdir(path=result\_path)

# Save the labels array to NPY file

np.save(

file=os.path.join(result\_path, "{}-gru\_svm-{}.npz".format(phase, step)),

arr=labels,

)