

```
# A Neural Network Architecture Combining Gated Recurrent Unit (GRU) and
# Support Vector Machine (SVM) for Intrusion Detection in Network Traffic Data
# Copyright (C) 2017 Abien Fred Agarap
#
# This program is free software: you can redistribute it and/or modify
# it under the terms of the GNU Affero General Public License as published
# by the Free Software Foundation, either version 3 of the License, or
# (at your option) any later version.
#
# This program is distributed in the hope that it will be useful,
# but WITHOUT ANY WARRANTY; without even the implied warranty of
# MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
# GNU Affero General Public License for more details.
#
# You should have received a copy of the GNU Affero General Public License
# along with this program. If not, see <http://www.gnu.org/licenses/>.
# =====

""""An implementation of L2-SVM model""""
from __future__ import absolute_import
from __future__ import division
from __future__ import print_function

__version__ = "0.3.7"
__author__ = "Abien Fred Agarap"

import numpy as np
import os
import sys
```

```
import tensorflow as tf
```

```
import time
```

```
class Svm:
```

```
    """Implementation of L2-Support Vector Machine using TensorFlow"""
```

```
    def __init__(self, alpha, batch_size, svm_c, num_classes, num_features):
```

```
        """Initialize the SVM class
```

```
        Parameter
```

```
        -----
```

```
        alpha : float
```

```
            The learning rate for the SVM model.
```

```
        batch_size : int
```

```
            Number of batches to use for training and testing.
```

```
        svm_c : float
```

```
            The SVM penalty parameter.
```

```
        num_classes : int
```

```
            Number of classes in a dataset.
```

```
        num_features : int
```

```
            Number of features in a dataset.
```

```
        """
```

```
        self.alpha = alpha
```

```
        self.batch_size = batch_size
```

```
        self.svm_c = svm_c
```

```
        self.num_classes = num_classes
```

```
        self.num_features = num_features
```

```

def __graph__():
    """Building the inference graph"""

    learning_rate = tf.placeholder(dtype=tf.float32, name="learning_rate")

    with tf.name_scope("input"):
        # [BATCH_SIZE, SEQUENCE_LENGTH]
        x_input = tf.placeholder(
            dtype=tf.float32, shape=[None, self.num_features], name="x_input"
        )

        # [BATCH_SIZE, N_CLASSES]
        y_input = tf.placeholder(dtype=tf.uint8, shape=[None], name="y_input")

        y_onehot = tf.one_hot(
            indices=y_input,
            depth=self.num_classes,
            on_value=1,
            off_value=-1,
            name="y_onehot",
        )

    with tf.name_scope("training_ops"):
        with tf.name_scope("weights"):
            weight = tf.get_variable(
                name="weights",
                initializer=tf.random_normal(
                    [self.num_features, self.num_classes], stddev=0.01
                ),
            ),

```

```

    )
    self.variable_summaries(weight)
with tf.name_scope("biases"):
    bias = tf.get_variable(
        name="biases",
        initializer=tf.constant(0.1, shape=[self.num_classes]),
    )
    self.variable_summaries(bias)
with tf.name_scope("Wx_plus_b"):
    y_hat = tf.matmul(x_input, weight) + bias
    tf.summary.histogram("pre-activations", y_hat)

# L2-SVM
with tf.name_scope("svm"):
    regularization = 0.5 * tf.reduce_sum(tf.square(weight))
    hinge_loss = tf.reduce_sum(
        tf.square(
            tf.maximum(
                tf.zeros([self.batch_size, self.num_classes]),
                1 - tf.cast(y_onehot, tf.float32) * y_hat,
            )
        )
    )
    with tf.name_scope("loss"):
        loss = regularization + self.svm_c * hinge_loss
tf.summary.scalar("loss", loss)

optimizer = tf.train.AdamOptimizer(learning_rate=learning_rate).minimize(
    loss

```

)

```
with tf.name_scope("accuracy"):
    predicted_class = tf.sign(y_hat)
    predicted_class = tf.identity(predicted_class, name="prediction")
    with tf.name_scope("correct_prediction"):
        correct = tf.equal(
            tf.argmax(predicted_class, 1), tf.argmax(y_onehot, 1)
        )
    with tf.name_scope("accuracy"):
        accuracy = tf.reduce_mean(tf.cast(correct, "float"))
tf.summary.scalar("accuracy", accuracy)
```

```
# merge all the summaries in the inference graph
merged = tf.summary.merge_all()
```

```
self.x_input = x_input
self.y_input = y_input
self.y_onehot = y_onehot
self.loss = loss
self.optimizer = optimizer
self.predicted_class = predicted_class
self.learning_rate = learning_rate
self.accuracy = accuracy
self.merged = merged
```

```
sys.stdout.write("\n<log> Building Graph...")
__graph__()
sys.stdout.write("</log>\n")
```

```
def train(
    self,
    checkpoint_path,
    log_path,
    model_name,
    epochs,
    result_path,
    train_data,
    train_size,
    validation_data,
    validation_size,
):
    """Trains the SVM model
```

Parameter

-----

checkpoint\_path : str

The directory where to save the trained model.

log\_path : str

The directory where to save the TensorBoard logs.

model\_name : str

The filename of the trained model.

epochs : int

The number of passes through the entire dataset.

result\_path : str

The path where to save the NPY files consisting of the actual and predicted labels.

train\_data : numpy.ndarray

The numpy.ndarray to be used as the training dataset.

train\_size : int

The number of data in `train\_data`.

validation\_data : numpy.ndarray

The numpy.ndarray to be used as the validation dataset.

validation\_size : int

The number of data in `validation\_data`.

"""

if not os.path.exists(checkpoint\_path):

os.mkdir(checkpoint\_path)

saver = tf.train.Saver(max\_to\_keep=1000)

# variable initializer

init\_op = tf.group(

tf.local\_variables\_initializer(), tf.global\_variables\_initializer()

)

# get the time tuple, and parse to str

timestamp = str(time.asctime())

# event file to contain TF graph summaries during training

train\_writer = tf.summary.FileWriter(

log\_path + timestamp + "-training", graph=tf.get\_default\_graph()

)

# event file to contain TF graph summaries during validation

validation\_writer = tf.summary.FileWriter(

log\_path + timestamp + "-validation", graph=tf.get\_default\_graph())

)

with tf.Session() as sess:

```
sess.run(init_op)
```

```
checkpoint = tf.train.get_checkpoint_state(checkpoint_path)
```

```
# check if a trained model exists
```

```
if checkpoint and checkpoint.model_checkpoint_path:
```

```
    # load the graph from the trained model
```

```
    saver = tf.train.import_meta_graph(  
        checkpoint.model_checkpoint_path + ".meta"
```

```
    )
```

```
# restore the variables
```

```
saver.restore(sess, tf.train.latest_checkpoint(checkpoint_path))
```

```
try:
```

```
    for step in range(epochs * train_size // self.batch_size):
```

```
        # set the value for slicing, to fetch batches of data
```

```
        offset = (step * self.batch_size) % train_size
```

```
        train_feature_batch = train_data[0][
```

```
            offset : (offset + self.batch_size)
```

```
        ]
```

```
        train_label_batch = train_data[1][
```

```
            offset : (offset + self.batch_size)
```

```
        ]
```



```

# dictionary for key-value pair input for training
feed_dict = {
    self.x_input: train_feature_batch,
    self.y_input: train_label_batch,
    self.learning_rate: self.alpha,
}

train_summary, _, predictions, actual = sess.run(
    [
        self.merged,
        self.optimizer,
        self.predicted_class,
        self.y_onehot,
    ],
    feed_dict=feed_dict,
)

# display training accuracy and loss every 100 steps and at step 0
if step % 100 == 0:

    # get the train loss and train accuracy
    train_accuracy, train_loss = sess.run(
        [self.accuracy, self.loss], feed_dict=feed_dict
    )

    # display the train loss and train accuracy
    print(
        "step [{}] train -- loss : {}, accuracy : {}".format(
            step, train_loss, train_accuracy

```

```

    )
)

# write the train summary
train_writer.add_summary(train_summary, step)

# save the model at the current time step
saver.save(sess, checkpoint_path + model_name, global_step=step)

self.save_labels(
    predictions=predictions,
    actual=actual,
    result_path=result_path,
    phase="training",
    step=step,
)
except KeyboardInterrupt:
    print("Training interrupted at {}".format(step))
    os._exit(1)
finally:
    print("EOF -- training done at step {}".format(step))

for step in range(epochs * validation_size // self.batch_size):

    offset = (step * self.batch_size) % validation_size
    validation_feature_batch = validation_data[0][
        offset : (offset + self.batch_size)
    ]
    validation_label_batch = validation_data[1][

```

```

        offset : (offset + self.batch_size)
    ]

    # dictionary for key-value pair input for validation
    feed_dict = {
        self.x_input: validation_feature_batch,
        self.y_input: validation_label_batch,
    }

    (
        test_summary,
        predictions,
        actual,
        test_loss,
        test_accuracy,
    ) = sess.run(
        [
            self.merged,
            self.predicted_class,
            self.y_onehot,
            self.loss,
            self.accuracy,
        ],
        feed_dict=feed_dict,
    )

    # display validation accuracy and loss every 100 steps
    if step % 100 == 0 and step > 0:

```

```

        # write the validation summary
        validation_writer.add_summary(test_summary, step)

    print(
        "step [{}] validation -- loss : {}, accuracy : {}".format(
            step, test_loss, test_accuracy
        )
    )

    self.save_labels(
        predictions=predictions,
        actual=actual,
        result_path=result_path,
        phase="validation",
        step=step,
    )

    print("EOF -- Testing done at step {}".format(step))

    @staticmethod
    def predict(
        batch_size, num_classes, test_data, test_size, checkpoint_path, result_path
    ):
        """Classifies the data whether there is an intrusion or none

        Parameter
        -----
        batch_size : int
            The number of batches to use for training/validation/testing.

```

num\_classes : int

The number of classes in a dataset.

test\_data : numpy.ndarray

The NumPy array testing dataset.

test\_size : int

The size of `test\_data`.

checkpoint\_path : str

The path where to save the trained model.

result\_path : str

The path where to save the actual and predicted classes array.

"""

# variables initializer

init\_op = tf.group(

tf.global\_variables\_initializer(), tf.local\_variables\_initializer()

)

with tf.Session() as sess:

sess.run(init\_op)

checkpoint = tf.train.get\_checkpoint\_state(checkpoint\_path)

# check if trained model exists

if checkpoint and checkpoint.model\_checkpoint\_path:

# load the trained model

saver = tf.train.import\_meta\_graph(

checkpoint.model\_checkpoint\_path + ".meta"

)

# restore the variables

```
saver.restore(sess, tf.train.latest_checkpoint(checkpoint_path))

print(
    "Loaded model from {}".format(
        tf.train.latest_checkpoint(checkpoint_path)
    )
)
```

try:

```
for step in range(test_size // batch_size):

    offset = (step * batch_size) % test_size

    test_example_batch = test_data[0][offset : (offset + batch_size)]
    test_label_batch = test_data[1][offset : (offset + batch_size)]

    # dictionary for input values for the tensors
    feed_dict = {"input/x_input:0": test_example_batch}

    # get the tensor for classification
    svm_tensor = sess.graph.get_tensor_by_name("accuracy/prediction:0")
    predictions = sess.run(svm_tensor, feed_dict=feed_dict)

    label_onehot = tf.one_hot(test_label_batch, num_classes, 1.0, -1.0)
    y_onehot = sess.run(label_onehot)

    # add key, value pair for labels
    feed_dict["input/y_input:0"] = test_label_batch

    # get the tensor for calculating the classification accuracy
    accuracy_tensor = sess.graph.get_tensor_by_name(
        "accuracy/accuracy/Mean:0"
```

```

    )

    accuracy = sess.run(accuracy_tensor, feed_dict=feed_dict)

    if step % 100 == 0 and step > 0:
        print("step [{}] test -- accuracy : {}".format(step, accuracy))

    Svm.save_labels(
        predictions=predictions,
        actual=y_onehot,
        result_path=result_path,
        step=step,
        phase="testing",
    )

except KeyboardInterrupt:
    print("KeyboardInterrupt at step {}".format(step))

finally:
    print("Done classifying at step {}".format(step))

```

@staticmethod

```

def variable_summaries(var):
    with tf.name_scope("summaries"):
        mean = tf.reduce_mean(var)
        tf.summary.scalar("mean", mean)
    with tf.name_scope("stddev"):
        stddev = tf.sqrt(tf.reduce_mean(tf.square(var - mean)))
    tf.summary.scalar("stddev", stddev)
    tf.summary.scalar("max", tf.reduce_max(var))
    tf.summary.scalar("min", tf.reduce_min(var))
    tf.summary.histogram("histogram", var)

```

@staticmethod

def save\_labels(predictions, actual, result\_path, step, phase):

"""Saves the actual and predicted labels to a NPY file

Parameter

-----

predictions : numpy.ndarray

The NumPy array containing the predicted labels.

actual : numpy.ndarray

The NumPy array containing the actual labels.

result\_path : str

The path where to save the concatenated actual and predicted labels.

step : int

The time step for the NumPy arrays.

phase : str

The phase for which the predictions is, i.e. training/validation/testing.

"""

# Concatenate the predicted and actual labels

labels = np.concatenate((predictions, actual), axis=1)

# Creates the result\_path directory if it does not exist

if not os.path.exists(path=result\_path):

os.mkdir(path=result\_path)

# save every labels array to NPY file

np.save(

file=os.path.join(result\_path, "{}-svm-{}.npz".format(phase, step)),



```
arr=labels,  
)
```